

**IMPLEMENTING OBSERVABILITY IN A SYSTEM USING  
PROMETHEUS AND GRAFANA**

by

MANPREET SINGH

Submitted in partial fulfilment of the requirements for the degree of

MASTER OF ENGINEERING

Major Subject: Internetworking

at

DALHOUSIE UNIVERSITY  
Halifax, Nova Scotia  
July, 2024

© Copyright by *Manpreet Singh*, 2024

**Dalhousie University**

**Faculty of Engineering**

**Internetworking**

The undersigned hereby certify that they have read and awarded a pass in INWK 6314 for the course/seminar entitled "IMPLEMENTING OBSERVABILITY IN A SYSTEM USING PROMETHEUS AND GRAFANA" by *Manpreet Singh* in partial fulfilment of the requirements for the degree of Master of Engineering.

---

**DALHOUSIE UNIVERSITY  
INTERNETWORKING PROGRAM**

**AUTHORITY TO DISTRIBUTE REPORT**

Title:

**IMPLEMENTING OBSERVABILITY IN A SYSTEM USING  
PROMETHEUS AND GRAFANA**

The Internetworking Program may make available or authorise others to make available individual photo/microfilm or soft copies of this report without restrictions after insert date here.

The author attests that permission has been obtained for the use of any copyrighted material appearing in this report (other than brief excerpts requiring only proper acknowledgement in scholarly writing) and that all such use is clearly acknowledged.

Full Name of Author:

Signature of Author: Manpreet Singh  
Date: July 5, 2024

## TABLE OF CONTENTS

<b>1 INTRODUCTION.....</b>	<b>1</b>
1.1 The Need For Observability.....	1
1.2 Implementing Observability On Localhost.....	2
<b>2 Components Of Prometheus.....</b>	<b>3</b>
2.1 Prometheus Server.....	3
2.2 PromQL.....	5
2.3 Time-Series Database.....	6
2.4 Exporters.....	6
2.5 Alertmanager.....	7
2.6 Integrating Grafana For Data Visualization.....	10
<b>3 FUTURE SCOPE OF THE PROJECT.....</b>	<b>11</b>
<b>4 REFERENCES.....</b>	<b>11</b>

## **ACKNOWLEDGEMENTS**

I would like to express my sincere gratitude to Dr. Hasan Raza for his invaluable guidance and support throughout the preparation of my project report. Your expertise and insightful feedback have been instrumental in shaping the document into its final form.

Your willingness to share your knowledge and provide constructive criticism have not only enhanced my understanding of the subject matter but have also significantly improved the quality of my work. Your mentorship has been a source of inspiration, and I am truly grateful for the time and effort you have invested in helping me succeed.

Thank you once again for your dedication and commitment to your students' academic growth. I look forward to applying the skills and knowledge I've gained under your guidance in my future endeavors.

## **EXECUTIVE SUMMARY**

Implementing observability into a system involves creating the capability to measure, monitor, and understand the internal state of the system. The implementation involves defining key metrics, storing the relevant information in a time series database, setting up dashboards for clarity in visualization(using Grafana) and Alert configurations using Prometheus Alertmanager based on defined Threshold/Metric parameters.

Metrics are numerical values that can be measured and tracked over time. They provide objective data points that reflect the state of the system.

Observability holds a significant space in the field of platform engineering as without observability, it is challenging to diagnose and troubleshoot issues quickly. When a problem occurs, it can be difficult to pinpoint the root cause, leading to increased downtime and frustration. Lack of metrics, logs, and traces means limited visibility into the system's behavior and performance, making it hard to understand what went wrong

The objective of the project is to delve into the workings of Prometheus and Grafana tools by scraping the metrics of the localhost where the system is deployed and generating alerts thereby notifying a threshold condition is reached and an appropriate action is required to bring the system back into optimal state.

# 1 INTRODUCTION

Prometheus is a robust open-source system and time-series database specifically designed for monitoring and observability. Developed initially by SoundCloud, Prometheus has emerged as a cornerstone in the landscape of cloud-native computing and DevOps practices. It excels in collecting metrics from diverse sources, empowering teams to monitor everything from infrastructure performance to application behavior in real-time. With its multidimensional data model, powerful query language (PromQL), and efficient storage, Prometheus enables proactive monitoring and troubleshooting across dynamic, distributed environments. Its ecosystem thrives on integrations with a wide array of software systems, fostering a community-driven approach that continuously evolves to meet the growing demands of scalability, reliability, and actionable insights in today's complex technological landscapes. Thus, Prometheus stands not only as a tool but as a symbol of empowerment for modern engineering teams striving for comprehensive observability and operational excellence.

## 1.1 The Need For Observability

Observability in a system refers to the ability to understand and diagnose what is happening inside the system based on its external outputs or behaviors. It is a critical aspect of system design and operation, particularly in complex and distributed systems. Here are several key reasons why observability is essential:

- **Diagnosing Issues:** In any system, problems can arise unexpectedly. Observability allows engineers and operators to monitor the system's behavior in real-time or retrospectively. This monitoring helps identify the root cause of issues quickly, leading to faster resolutions and minimizing downtime.
- **Understanding Performance:** Observability provides insights into how the system is performing under normal conditions and during peak loads or stress. This visibility helps in optimizing performance and resource allocation, ensuring efficient use of system resources.
- **Monitoring User Experience:** For systems that interact directly with users (such as web applications or customer service platforms), observability helps track and understand user interactions and experiences. This information is crucial for improving user satisfaction and addressing usability issues.

- Analyzing Trends and Patterns: By collecting and analyzing observational data over time, patterns and trends can be identified. These insights enable proactive decision-making, such as capacity planning, feature enhancements, or detecting potential security threats.
- Debugging and Iterative Development: During the development phase, observability tools aid in debugging code and verifying new features. Developers can monitor how changes affect the system and validate performance against expected outcomes.

## 1.2 Implementing Observability On Localhost

In order to execute the project, we have deployed a stack of services on the local host. The official documentation(<https://prometheus.io/>) was referenced extensively to understand the depth of the system.

## 2 COMPONENTS OF PROMETHEUS

The emphasis of this section is to understand the core components of the Prometheus tool. We have divided this section into subsections to describe the different components Prometheus provides to implement observability.

## 2.1 Prometheus Server

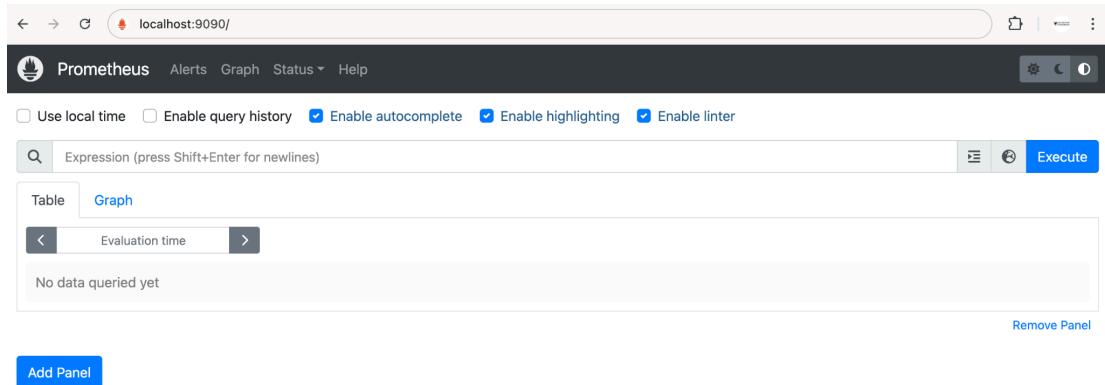
Prometheus Server stands as the cornerstone of the Prometheus observability ecosystem, pivotal in modern cloud-native monitoring and alerting practices. Functioning as both collector and storage engine, it efficiently gathers time-series metrics from diverse targets via HTTP endpoints, employing a multi-dimensional data model with labels for flexible querying and aggregation. Supported by PromQL, its powerful query language, Prometheus Server enables detailed analysis and real-time monitoring of application and infrastructure performance. With robust alerting rules and integration with Alertmanager for streamlined alert handling, it empowers proactive incident response. Its scalable architecture, coupled with support for various service discovery mechanisms and exporters, ensures adaptability to dynamic environments.

Prometheus Server provides a built-in web-based user interface (UI) for exploring metrics, running queries, and visualizing data. Additionally, it exposes a comprehensive HTTP API that allows programmatic access for querying metrics, managing alerting rules, and more.

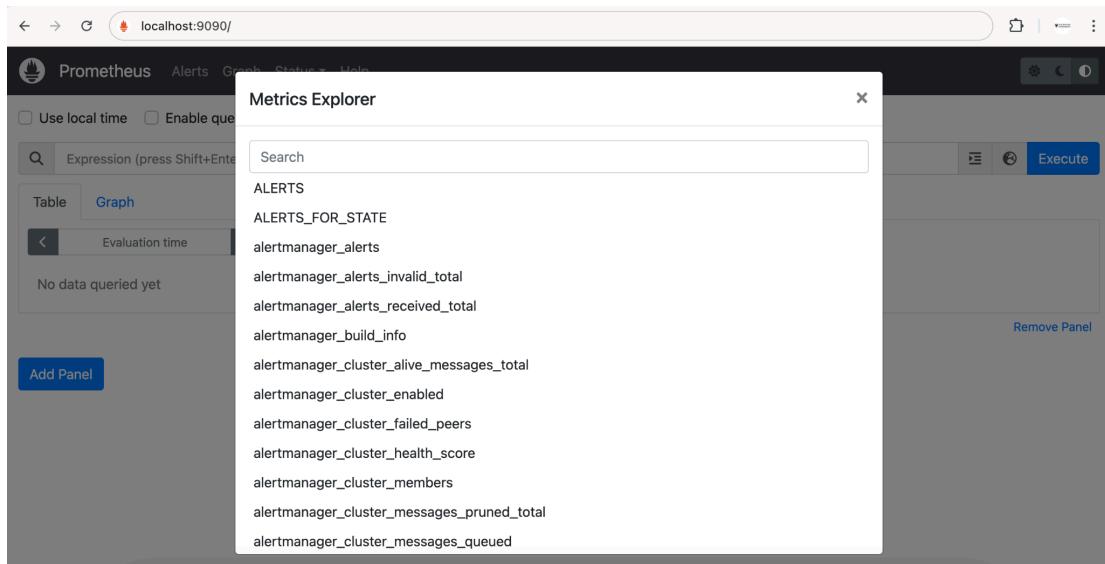
The Prometheus server runs on port 9090.

### Prometheus service running on localhost:

## Prometheus UI:



## Metrics Exposed From Prometheus:



## Prometheus Server Configuration File:

```

Prometheus Project > Prometheus Server > prometheus-2.53.0.darwin-amd64 > prometheus.yml
Project -> prometheus.yml

1   global:
2     scrape_interval: 15s # Set the scrape interval to every 15 seconds. Default is every 1 minute.
3     evaluation_interval: 15s # Evaluate rules every 15 seconds. The default is every 1 minute.
4   alerting:
5     alertmanagers:
6       - static_configs:
7         - targets:
8           - 'localhost:9093'
9   rule_files:
10    - "rules/alert.yml"
11   scrape_configs:
12     - job_name: "prometheus"
13       static_configs:
14         - targets: ["localhost:9090"]
15     - job_name: "node_exporter"
16       static_configs:
17         - targets: ["localhost:9100"]
18     - job_name: "alertmanager"
19       static_configs:
20         - targets: ["localhost:9093"]

```

## 2.2 PromQL

PromQL or Prometheus Query Language is a powerful query language designed for working with Prometheus data. It allows users to write expressive queries to retrieve, manipulate, and aggregate time-series data based on specific criteria

It is important to reveal all the details of the problem to the reader. This chapter is used to make the reader aware of all the nuances of the problem so that any subsequent trade-offs or measurements are justifiable in the context of the problem desc

## 2.3 Time-Series Database

Prometheus Server maintains a local, efficient time-series database optimized for storing metrics collected from various targets. This database structure allows for high write throughput and efficient querying of historical data

## 2.4 Exporters

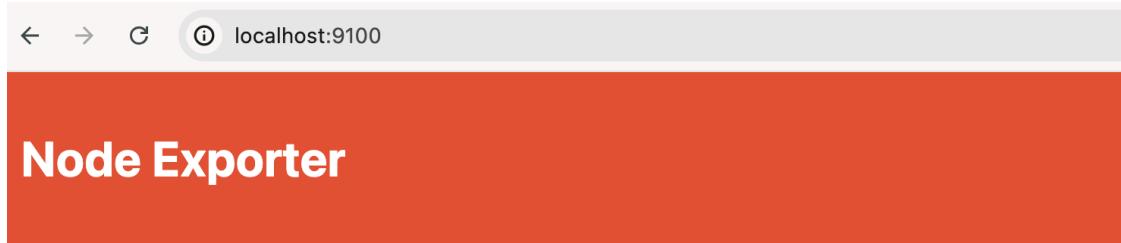
Exporters in Prometheus play a pivotal role in extending the monitoring capabilities of the Prometheus ecosystem by enabling the collection of metrics from a diverse array of systems and services that do not natively expose Prometheus-compatible metrics. These exporters act as intermediaries, translating metrics from various formats or protocols into a format Prometheus can scrape and store in its time-series database. They facilitate seamless integration with a wide range of systems, including databases (such as MySQL, PostgreSQL), web servers (like Apache and Nginx), message brokers (such as RabbitMQ and Kafka), cloud platforms (like AWS and Google Cloud), and numerous other applications and services crucial to modern IT infrastructures. Each exporter is specifically designed to understand the metrics format of the system it monitors and exposes these metrics via an HTTP endpoint typically located at `/metrics`. Prometheus scrapes these endpoints at configured intervals, allowing it to collect metrics on resource utilization, performance indicators, operational metrics, and more. Popular examples of exporters include the Node Exporter for system metrics like CPU, memory, and disk usage; the Blackbox Exporter for probing endpoints via HTTP, ICMP, TCP, and DNS; and the JMX

We have considered the use case of deploying NodeExporter on local host which runs the service on port 9100 and also exposes a UI on <http://localhost:9100>

### NodeExporter Service Deployed On Localhost:

```
manpreetsingh@tc69f node_exporter-1.8.1.darwin-amd64 % ./node_exporter
ts=2024-07-03T23:19:37.751Z caller=node_exporter.go:193 level=info msg="Starting node_exporter" version="(version=1.8.1, branch=HEAD, revision=400c3
979931613db930ea95f39cecb377cdbb5b)"
ts=2024-07-03T23:19:37.755Z caller=node_exporter.go:194 level=info msg="Build context" build_context="(go=go1.22.3, platform=darwin/amd64, user=root
@eadd40c03e38, date=20240521-18:39:19, tags=unknown)"
ts=2024-07-03T23:19:37.756Z caller=filesystem_common.go:111 level=info collector=filesystem msg="Parsed flag --collector.filesystem.mount-points-exclu
de" flag=/^(dev|$/)
ts=2024-07-03T23:19:37.757Z caller=filesystem_common.go:113 level=info collector=filesystem msg="Parsed flag --collector.filesystem.fs-types-exclude
" flag=~/devfs$
ts=2024-07-03T23:19:37.757Z caller=node_exporter.go:111 level=info msg="Enabled collectors"
ts=2024-07-03T23:19:37.757Z caller=node_exporter.go:118 level=info collector=boottime
ts=2024-07-03T23:19:37.757Z caller=node_exporter.go:118 level=info collector=cpu
ts=2024-07-03T23:19:37.757Z caller=node_exporter.go:118 level=info collector=disksstats
ts=2024-07-03T23:19:37.757Z caller=node_exporter.go:118 level=info collector=filesystem
ts=2024-07-03T23:19:37.757Z caller=node_exporter.go:118 level=info collector=loadavg
ts=2024-07-03T23:19:37.757Z caller=node_exporter.go:118 level=info collector=meminfo
ts=2024-07-03T23:19:37.757Z caller=node_exporter.go:118 level=info collector=netdev
ts=2024-07-03T23:19:37.757Z caller=node_exporter.go:118 level=info collector=oss
ts=2024-07-03T23:19:37.757Z caller=node_exporter.go:118 level=info collector=powersupplyclass
ts=2024-07-03T23:19:37.757Z caller=node_exporter.go:118 level=info collector=textfile
ts=2024-07-03T23:19:37.757Z caller=node_exporter.go:118 level=info collector=thermal
ts=2024-07-03T23:19:37.757Z caller=node_exporter.go:118 level=info collector=time
ts=2024-07-03T23:19:37.757Z caller=node_exporter.go:118 level=info collector=uname
ts=2024-07-03T23:19:37.776Z caller=tls_config.go:313 level=info msg="Listening on" address=[::]:9100
ts=2024-07-03T23:19:37.776Z caller=tls_config.go:316 level=info msg="TLS is disabled." http2=false address=[::]:9100
```

### NodeExporter UI Exposed On Port 9100:



### Prometheus Node Exporter

Version: (version=1.8.1, branch=HEAD, revision=400c3979931613db930ea035f39ce7b377cdbb5b)

- [Metrics](#)

'Metrics' would display a list of all the metrics which are exposed from the NodeExporter thereby exposing time-series data points of various metrics of the localhost.

## 2.5 Alertmanager

Prometheus Server supports alerting based on predefined rules and thresholds. When conditions specified in alerting rules are met, Prometheus triggers alerts. These alerts are then managed and routed by Alertmanager, a separate component closely integrated with Prometheus Server. This component runs on port 9093 and has its own configuration file to configure various alerts and send the alerts to an external system.

**Keeping the scope of the project in mind, I have configured gmail as the external system receiving the alerts when localhost's CPU utilization crosses 20% utilization.**

## AlertManager Running On Port 9093 on Localhost:

```
manpreetsingh@Manpreets-MBP alertmanager -v=0.27.0 darwin-amd64 %
manpreetsingh@Manpreets-MBP alertmanager -v=0.27.0 darwin-amd64 % ./alertmanager
ts=2024-07-03T02:33:22.897Z caller=main.go:181 level=info msg="Starting Alertmanager" version="(version=0.27.0, branch=HEAD, revision=0aa3c2aad14cff039931923ab16b26b7481783b5)" ts=2024-07-03T02:33:22.897Z caller=main.go:182 level=info build_context="(go=g01.21.7, platform=darwin/amd64, user=root@cb7255915c77, date=20240228-11:47:50, tags=netgo)"
ts=2024-07-03T02:33:22.913Z caller=cluster.go:186 level=info component=cluster msg="setting advertise address explicitly" addr=192.168.4.91 port=9094
ts=2024-07-03T02:33:22.926Z caller=cluster.go:683 level=info component=cluster msg="Waiting for gossip to settle..." interval=2s
ts=2024-07-03T02:33:22.973Z caller=coordinator.go:113 level=info component=configuration msg="Loading configuration file" file=alertmanager.yml
ts=2024-07-03T02:33:22.980Z caller=coordinator.go:126 level=info components=configuration msg="Completed loading of configuration file" file=alertmanager.yml
ts=2024-07-03T02:33:22.981Z caller=coordinator.go:128 level=info component=cluster msg="Setting cluster address[:] to 19093"
ts=2024-07-03T02:33:22.982Z caller=tls_config.go:314 level=info msg="TLS is disabled." http2=false address[:]::9093
ts=2024-07-03T02:33:22.982Z caller=cluster.go:788 level=info component=cluster msg="gossip not settled" polls=0 before=0 now=1 elapsed=2.000546243s
ts=2024-07-03T02:33:22.927Z caller=cluster.go:788 level=info component=cluster msg="gossip settled; proceeding" elapsed=10.007156277s
^Cts=2024-07-03T16:41:41.255Z caller=main.go:542 level=info msg="Received SIGTERM, exiting gracefully..."
manpreetsingh@Manpreets-MBP alertmanager -v=0.27.0 darwin-amd64 % ./alertmanager
ts=2024-07-03T16:41:42.392Z caller=main.go:181 level=info msg="Starting Alertmanager" version="(version=0.27.0, branch=HEAD, revision=0aa3c2aad14cff039931923ab16b26b7481783b5)" ts=2024-07-03T16:41:42.392Z caller=main.go:182 level=info build_context="(go=g01.21.7, platform=darwin/amd64, user=root@cb7255915c77, date=20240228-11:47:50, tags=netgo)"
ts=2024-07-03T16:41:42.392Z caller=cluster.go:113 level=info component=cluster msg="Setting cluster address[:] to 19093"
ts=2024-07-03T16:41:42.392Z caller=coordinator.go:113 level=info component=configuration msg="Completed loading of configuration file" file=alertmanager.yml
ts=2024-07-03T16:41:42.392Z caller=coordinator.go:126 level=info component=configuration msg="Completed loading of configuration file" file=alertmanager.yml
ts=2024-07-03T16:41:42.392Z caller=tls_config.go:313 level=info msg="Listening on" address[:]::9093
ts=2024-07-03T16:41:42.392Z caller=cluster.go:788 level=info component=cluster msg="gossip not settled" polls=0 before=0 now=1 elapsed=2.0008754005s
ts=2024-07-03T16:41:42.392Z caller=cluster.go:788 level=info component=cluster msg="gossip settled; proceeding" elapsed=10.007076644s
^Cts=2024-07-03T16:41:42.392Z caller=main.go:542 level=info msg="Received SIGTERM, exiting gracefully..."
manpreetsingh@Manpreets-MBP alertmanager -v=0.27.0 darwin-amd64 % ./alertmanager
ts=2024-07-03T16:51:15.876Z caller=main.go:181 level=info msg="Starting Alertmanager" version="(version=0.27.0, branch=HEAD, revision=0aa3c2aad14cff039931923ab16b26b7481783b5)" ts=2024-07-03T16:51:15.876Z caller=main.go:182 level=info build_context="(go=g01.21.7, platform=darwin/amd64, user=root@cb7255915c77, date=20240228-11:47:50, tags=netgo)"
ts=2024-07-03T16:51:15.886Z caller=cluster.go:113 level=info component=cluster msg="Setting cluster address[:] to 19094"
ts=2024-07-03T16:51:15.888Z caller=coordinator.go:113 level=info component=configuration msg="Completed loading of configuration file" file=alertmanager.yml
ts=2024-07-03T16:51:15.888Z caller=coordinator.go:126 level=info component=configuration msg="Completed loading of configuration file" file=alertmanager.yml
ts=2024-07-03T16:51:15.888Z caller=tls_config.go:313 level=info msg="Listening on" address[:]::9094
ts=2024-07-03T16:51:15.892Z caller=cluster.go:788 level=info component=cluster msg="gossip not settled" polls=0 before=0 now=1 elapsed=2.0002616762s
ts=2024-07-03T16:51:15.892Z caller=cluster.go:788 level=info component=cluster msg="gossip settled; proceeding" elapsed=10.005973832s
^Cts=2024-07-03T16:51:15.892Z caller=main.go:542 level=info msg="Received SIGTERM, exiting gracefully..."
manpreetsingh@Manpreets-MBP alertmanager -v=0.27.0 darwin-amd64 % ./alertmanager
ts=2024-07-03T16:51:15.892Z caller=main.go:181 level=info msg="Starting Alertmanager" version="(version=0.27.0, branch=HEAD, revision=0aa3c2aad14cff039931923ab16b26b7481783b5)" ts=2024-07-03T16:51:15.892Z caller=main.go:182 level=info build_context="(go=g01.21.7, platform=darwin/amd64, user=root@cb7255915c77, date=20240228-11:47:50, tags=netgo)"
ts=2024-07-03T16:51:15.892Z caller=cluster.go:113 level=info component=cluster msg="Setting cluster address[:] to 19094"
ts=2024-07-03T16:51:15.892Z caller=coordinator.go:113 level=info component=configuration msg="Completed loading of configuration file" file=alertmanager.yml
ts=2024-07-03T16:51:15.892Z caller=coordinator.go:126 level=info component=configuration msg="Completed loading of configuration file" file=alertmanager.yml
ts=2024-07-03T16:51:15.892Z caller=tls_config.go:313 level=info msg="Listening on" address[:]::9094
ts=2024-07-03T16:51:15.892Z caller=cluster.go:788 level=info component=cluster msg="gossip not settled" polls=0 before=0 now=1 elapsed=2.0002616762s
ts=2024-07-03T16:51:15.892Z caller=cluster.go:788 level=info component=cluster msg="gossip settled; proceeding" elapsed=10.005973832s
^Cts=2024-07-03T16:51:15.892Z caller=main.go:542 level=info msg="Received SIGTERM, exiting gracefully..."
manpreetsingh@Manpreets-MBP alertmanager -v=0.27.0 darwin-amd64 % ./alertmanager
ts=2024-07-03T17:00:11.478Z caller=main.go:182 level=info build_context="(go=g01.21.7, platform=darwin/amd64, user=root@cb7255915c77, date=20240228-11:47:50, tags=netgo)"
ts=2024-07-03T17:00:11.493Z caller=cluster.go:186 level=info component=cluster msg="setting advertise address explicitly" addr=192.168.4.91 port=9094
```

## AlertManagerUI:

The screenshot shows the AlertManager UI interface. At the top, there's a header with a back arrow, a forward arrow, and a refresh icon. To the right of these is the URL "localhost:9093/#/status". Below the header is a navigation bar with tabs: "Alertmanager", "Alerts", "Silences", "Status", "Settings", and "Help". The "Status" tab is currently active. The main content area has a title "Status" and a timestamp "Uptime: 2024-07-03T18:13:54.537Z". Below this, there's a section titled "Cluster Status" with a "Name" field containing "01J1WWV57746Q5T5RY43Q2Z1GB" and a "Status" field showing "ready". Under "Peers", there's a list item with a "Name" of "01J1WWV57746Q5T5RY43Q2Z1GB" and an "Address" of "192.168.4.91:9094".

## Status

**Uptime:** 2024-07-03T18:13:54.537Z

## Cluster Status

**Name:** 01J1WWV57746Q5T5RY43Q2Z1GB

**Status:** ready

**Peers:**

- Name:** 01J1WWV57746Q5T5RY43Q2Z1GB
- Address:** 192.168.4.91:9094

## Version Information

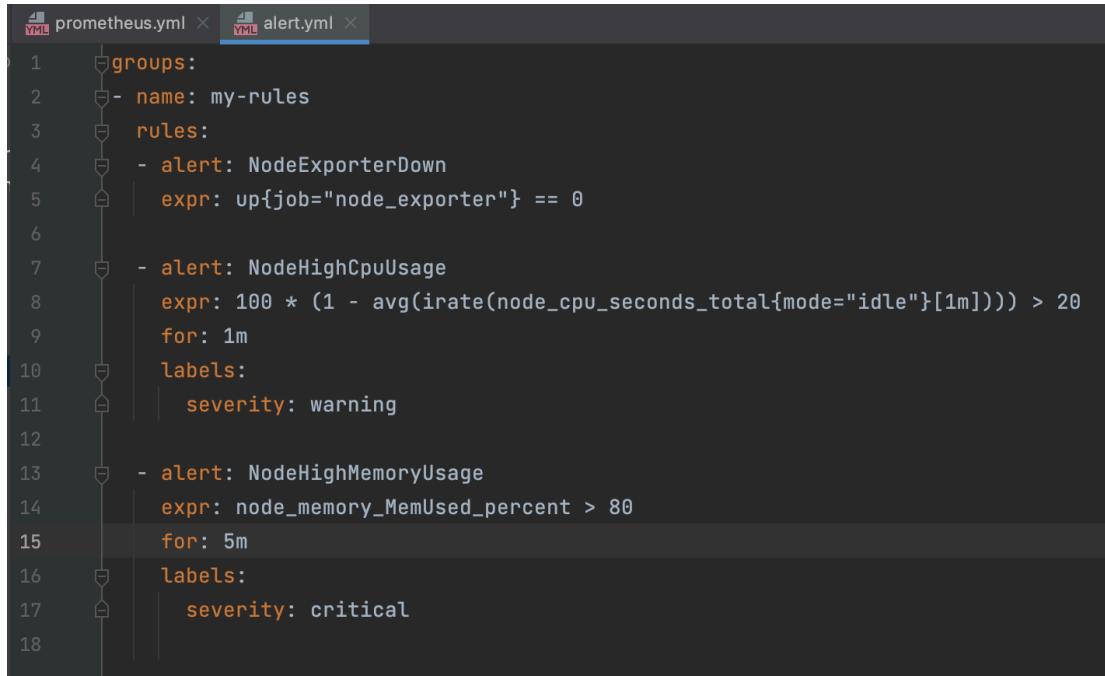
**Branch:** HEAD

**BuildDate:** 20240228-11:47:50

**BuildUser:** root@cb7255915c77

**GoVersion:** go1.21.7

### Alerts Configured Referring To NodeExporter Metrics:



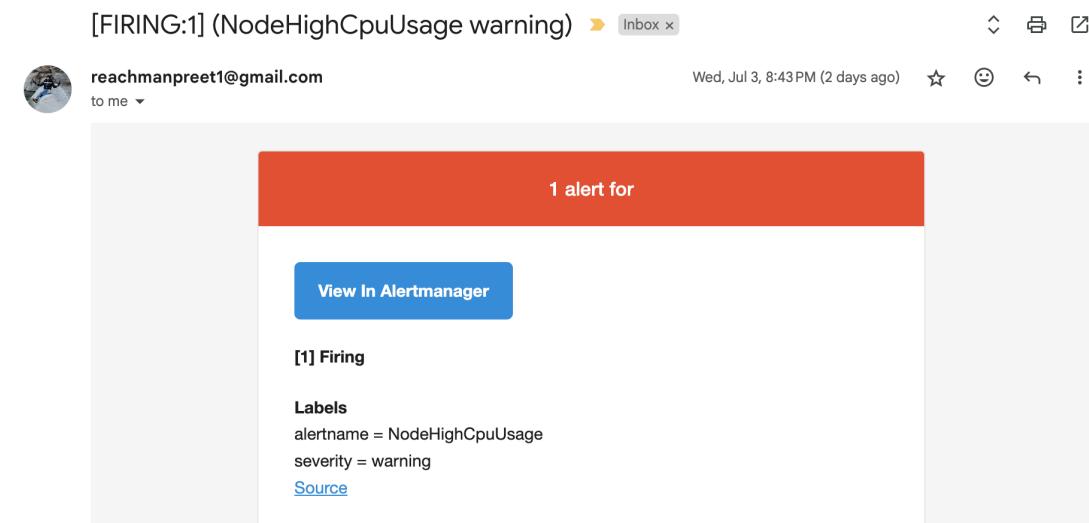
```

1   groups:
2     - name: my-rules
3       rules:
4         - alert: NodeExporterDown
5           expr: up{job="node_exporter"} == 0
6
7         - alert: NodeHighCpuUsage
8           expr: 100 * (1 - avg(irate(node_cpu_seconds_total{mode="idle"}[1m]))) > 20
9           for: 1m
10          labels:
11            severity: warning
12
13        - alert: NodeHighMemoryUsage
14          expr: node_memory_MemUsed_percent > 80
15          for: 5m
16          labels:
17            severity: critical
18

```

This alert.yml file is then pointed into the Prometheus.yaml file so prometheus can read the alerts and fire them to AlertManager

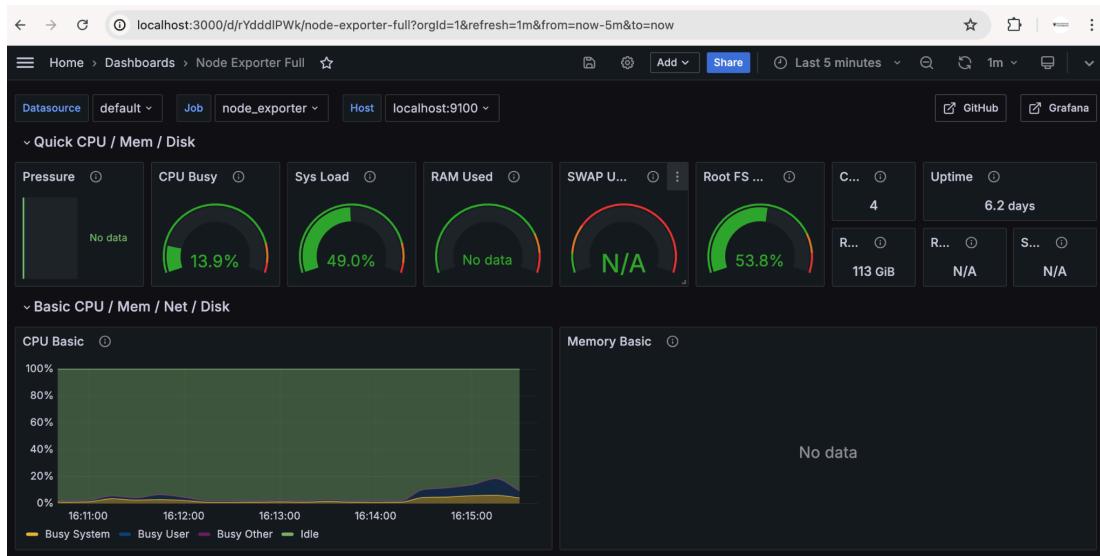
### Alerts Triggered To The Gmail System From AlertManager:



## 2.6 Integrating Grafana For Data Visualization

We have imported the NodeExporter and Prometheus templates in our project in the Grafana tool which runs on port 3000 on the localhost.

## Grafana Tool Exposing The NodeExporter Metrics:



## Grafana Service Running On Port 3000 On LocalHost:

### 3 FUTURE SCOPE OF THE PROJECT

Since Prometheus, AlertManager and Grafana are not only used to observe metrics of Infrastructure, the next level of implementation would be to also consider instrumenting application code and scrape application specific metrics and widen the scope of infrastructure to, for example, the Kubernetes platform as Prometheus works natively with the Kubernetes deployed infrastructure and execute observability on tier 0 (highly critical) applications and infrastructure to use this platform to its most effective level.

### 4 REFERENCES

- <https://prometheus.io/>
- <https://www.udemy.com/course/prometheus-course/>
- <https://techcrunch.com/>